

Problem Set2: Wrangling Subway Data

1 - Number of Rainy Days:

```
import pandas
```

```
import pandasql
```

```
def num_rainy_days(filename):
```

```
'''
```

This function should run a SQL query on a dataframe of weather data. The SQL query should return one column and one row - a count of the number of days in the dataframe where the rain column is equal to 1 (i.e., the number of days it rained). The dataframe will be titled 'weather_data'. You'll need to provide the SQL query. You might find SQL's count function useful for this exercise. You can read more about it here:

<https://dev.mysql.com/doc/refman/5.1/en/counting-rows.html>

You might also find that interpreting numbers as integers or floats may not work initially. In order to get around this issue, it may be useful to cast these numbers as integers. This can be done by writing `cast(column as integer)`. So for example, if we wanted to cast the `maxtempi` column as an integer, we would actually write something like `where cast(maxtempi as integer) = 76`, as opposed to simply `where maxtempi = 76`.

You can see the weather data that we are passing in below:

https://www.dropbox.com/s/7sf0yqc9ykpq3w8/weather_underground.csv

```
'''
```

```
weather_data = pandas.read_csv(filename)
```

```
q = """
```

```
SELECT COUNT(*) FROM weather_data WHERE CAST(rain as integer) = 1
```

```
"""
```

```
#Execute your SQL command against the pandas frame
```

```
rainy_days = pandasql.sqldf(q.lower(), locals())
```

```
return rainy_days
```

```
count(*)
```

```
0          10
```

2 - Temp on Foggy and Nonfoggy Days:

```
import pandas
```

```
import pandasql
```

```
def max_temp_aggregate_by_fog(filename):
```

```
'''
```

This function should run a SQL query on a dataframe of weather data. The SQL query should return two columns and two rows - whether it was foggy or not (0 or 1) and the max maxtempi for that fog value (i.e., the maximum max temperature for both foggy and non-foggy days). The dataframe will be titled 'weather_data'. You'll need to provide the SQL query.

You might also find that interpreting numbers as integers or floats may not work initially. In order to get around this issue, it may be useful to cast these numbers as integers. This can be done by writing `cast(column as integer)`. So for example, if we wanted to cast the maxtempi column as an integer, we would actually write something like `where cast(maxtempi as integer) = 76`, as opposed to simply `where maxtempi = 76`.

You can see the weather data that we are passing in below:

https://www.dropbox.com/s/7sf0yqc9ykpq3w8/weather_underground.csv

```
'''
```

```
weather_data = pandas.read_csv(filename)
```

```
q = """
```

```
SELECT fog, max(CAST (maxtempi as integer)) FROM weather_data GROUP BY fog
```

```
"""
```

```
#Execute your SQL command against the pandas frame
```

```
foggy_days = pandasql.sqldf(q.lower(), locals())
```

```
return foggy_days
```

Output by your program below.

	fog	max(cast (maxtempi as integer))
0	0	86
1	1	81

3 - Mean Temp on Weekends:

import pandas

import pandasql

def avg_weekend_temperature(filename):

'''

This function should run a SQL query on a dataframe of weather data. The SQL query should return one column and one row - the average meantempi on days that are a Saturday or Sunday (i.e., the the average mean temperature on weekends). The dataframe will be titled 'weather_data' and you can access the date in the dataframe via the 'date' column.

You'll need to provide the SQL query.

You might also find that interpreting numbers as integers or floats may not work initially. In order to get around this issue, it may be useful to cast these numbers as integers. This can be done by writing `cast(column as integer)`. So for example, if we wanted to cast the `maxtempi` column as an integer, we would actually write something like `where cast(maxtempi as integer) = 76`, as opposed to simply `where maxtempi = 76`.

Also, you can convert dates to days of the week via the `'strftime'` keyword in SQL. For example, `cast (strftime('%w', date) as integer)` will return 0 if the date is a Sunday or 6 if the date is a Saturday.

You can see the weather data that we are passing in below:

https://www.dropbox.com/s/7sf0yqc9ykpq3w8/weather_underground.csv

'''

weather_data = pandas.read_csv(filename)

q = """

SELECT avg(cast (meantempi as integer)) FROM weather_data WHERE CAST(strftime('%w', date) as integer) = 0 or CAST(strftime('%w', date) as integer) = 6

"""

#Execute your SQL command against the pandas frame

mean_temp_weekends = pandasql.sqldf(q.lower(), locals())

return mean_temp_weekends

Output by your program below.

```
avg(cast (meantempi as integer))
0                                65.111111
```

4 - Mean Temp on Rainy Days:

import pandas

import pandasql

def avg_min_temperature(filename):

'''

This function should run a SQL query on a dataframe of weather data. More specifically you want to find the average minimum temperature on rainy days where the minimum temperature is greater than 55 degrees.

You might also find that interpreting numbers as integers or floats may not work initially. In order to get around this issue, it may be useful to cast these numbers as integers. This can be done by writing `cast(column as integer)`. So for example, if we wanted to cast the `maxtempi` column as an integer, we would actually write something like `where cast(maxtempi as integer) = 76`, as opposed to simply `where maxtempi = 76`.

You can see the weather data that we are passing in below:

https://www.dropbox.com/s/7sf0yqc9ykpq3w8/weather_underground.csv

'''

weather_data = pandas.read_csv(filename)

q = """

SELECT avg(CAST (mintempi as integer)) FROM weather_data WHERE rain = 1 AND mintempi > 55

"""

#Execute your SQL command against the pandas frame

avg_min_temp_rainy = pandasql.sqldf(q.lower(), locals())

return avg_min_temp_rainy

Output by your program below.

```
avg(cast (mintempi as integer))
0                                61.25
```

5 - Fixing Turnstile Data:

import csv

def fix_turnstile_data(filenames):

'''

Filenames is a list of MTA Subway turnstile text files. A link to an example MTA Subway turnstile text file can be seen at the URL below:

http://web.mta.info/developers/data/nyct/turnstile/turnstile_110507.txt

As you can see, there are numerous data points included in each row of the a MTA Subway turnstile text file.

You want to write a function that will update each row in the text file so there is only one entry per row. A few examples below:

A002,R051,02-00-00,05-28-11,00:00:00,REGULAR,003178521,001100739

A002,R051,02-00-00,05-28-11,04:00:00,REGULAR,003178541,001100746

A002,R051,02-00-00,05-28-11,08:00:00,REGULAR,003178559,001100775

Write the updates to a different text file in the format of "updated_" + filename.

For example:

1) if you read in a text file called "turnstile_110521.txt"

2) you should write the updated data to "updated_turnstile_110521.txt"

The order of the fields should be preserved. Remember to read through the Instructor Notes below for more details on the task.

In addition, here is a CSV reader/writer introductory tutorial:

<http://goo.gl/HBbvyy>

You can see a sample of the turnstile text file that's passed into this function and the the corresponding updated file in the links below:

Sample input file:

https://www.dropbox.com/s/mpin5zv4hgrx244/turnstile_110528.txt

Sample updated file:

https://www.dropbox.com/s/074xbgio4c39b7h/solution_turnstile_110528.txt

'''

for name in filenames:

f_in = open(name, 'r')

f_out = open("updated_" + name, 'w')

```

reader_in = csv.reader(f_in)
reader_out = csv.writer(f_out)

for line in reader_in:
    index = 3
    header=line[0:3]
    length = len(line)
    length2 = length - 3

    for index in range(3, length2 + 1):
        fw = header + line[index:(index+5)]
        reader_out.writerow(fw)
        index = index + 5

return filenames

```

Good job. Your code worked perfectly.
Your code produced the following output:

updated_turnstile_110528.txt

6 - Combining Turnstile Data:

```

def create_master_turnstile_file(filenames, output_file):
'''

```

Write a function that takes the files in the list filenames, which all have the columns 'C/A, UNIT, SCP, DATEn, TIMEn, DESCn, ENTRIESn, EXITSn', and consolidates them into one file located at output_file. There should be ONE row with the column headers, located at the top of the file. The input files do not have column header rows of their own.

For example, if file_1 has:

'C/A, UNIT, SCP, DATEn, TIMEn, DESCn, ENTRIESn, EXITSn'

line 1 ...

line 2 ...

and another file, file_2 has:

'C/A, UNIT, SCP, DATEn, TIMEn, DESCn, ENTRIESn, EXITSn'

line 3 ...

line 4 ...

line 5 ...

We need to combine file_1 and file_2 into a master_file like below:

```
'C/A, UNIT, SCP, DATEn, TIMEn, DESCn, ENTRIESn, EXITSn'
```

line 1 ...

line 2 ...

line 3 ...

line 4 ...

line 5 ...

```
'''
```

with open(output_file, 'w') as master_file:

```
    master_file.write('C/A,UNIT,SCP,DATEn,TIMEn,DESCn,ENTRIESn,EXITSn\n')
```

for filename in filenames:

```
    # your code here
```

with open(filename) as infile:

```
    for line in infile:
```

```
        master_file.write(line)
```

Good job. Your code worked perfectly.

Your code produced the following output:

```
C/A,UNIT,SCP,DATEn,TIMEn,DESCn,ENTRIESn,EXITSn
```

7 - Filtering Irregular Data:

import pandas

def filter_by_regular(filename):

```
'''
```

This function should read the csv file located at filename into a pandas dataframe, and filter the dataframe to only rows where the 'DESCn' column has the value 'REGULAR'.

For example, if the pandas dataframe is as follows:

```
,C/A,UNIT,SCP,DATEn,TIMEn,DESCn,ENTRIESn,EXITSn
```

```
0,A002,R051,02-00-00,05-01-11,00:00:00,REGULAR,3144312,1088151
```

```
1,A002,R051,02-00-00,05-01-11,04:00:00,DOOR,3144335,1088159
```

```
2,A002,R051,02-00-00,05-01-11,08:00:00,REGULAR,3144353,1088177
```

```
3,A002,R051,02-00-00,05-01-11,12:00:00,DOOR,3144424,1088231
```

The dataframe will look like below after filtering to only rows where DESCn column has the value 'REGULAR':

```
0,A002,R051,02-00-00,05-01-11,00:00:00,REGULAR,3144312,1088151
2,A002,R051,02-00-00,05-01-11,08:00:00,REGULAR,3144353,1088177
'''
```

```
turnstile_data = pandas.read_csv(filename)
# your code here
# more of your code here
turnstile_data = pandas.DataFrame(turnstile_data)
turnstile_data = turnstile_data[(turnstile_data.DEScN == 'REGULAR')]

return turnstile_data
```

Good job! Your code worked perfectly. Your output below:

	C/A	UNIT	SCP	DATEn	TIMEn	DESCn	ENTRIESn	EXITSn
0	A002	R051	02-00-00	05-01-11	00:00:00	REGULAR	3144312	1088151

8 - Get Hourly Entries:

```
import pandas
```

```
def get_hourly_entries(df):
```

```
'''
```

The data in the MTA Subway Turnstile data reports on the cumulative number of entries and exits per row. Assume that you have a dataframe called df that contains only the rows for a particular turnstile machine (i.e., unique SCP, C/A, and UNIT). This function should change these cumulative entry numbers to a count of entries since the last reading (i.e., entries since the last row in the dataframe).

More specifically, you want to do two things:

- 1) Create a new column called ENTRIESn_hourly
- 2) Assign to the column the difference between ENTRIESn of the current row and the previous row. If there is any NaN, fill/replace it with 1.

You may find the pandas functions shift() and fillna() to be helpful in this exercise.

Examples of what your dataframe should look like at the end of this exercise:

	C/A	UNIT	SCP	DATE	TIME	DESC	ENTRIES	EXIT	ENTRIES_hourly
0	A002	R051	02-00-00	05-01-11	00:00:00	REGULAR	3144312	1088151	1
1	A002	R051	02-00-00	05-01-11	04:00:00	REGULAR	3144335	1088159	23
2	A002	R051	02-00-00	05-01-11	08:00:00	REGULAR	3144353	1088177	18
3	A002	R051	02-00-00	05-01-11	12:00:00	REGULAR	3144424	1088231	71
4	A002	R051	02-00-00	05-01-11	16:00:00	REGULAR	3144594	1088275	170
5	A002	R051	02-00-00	05-01-11	20:00:00	REGULAR	3144808	1088317	214
6	A002	R051	02-00-00	05-02-11	00:00:00	REGULAR	3144895	1088328	87
7	A002	R051	02-00-00	05-02-11	04:00:00	REGULAR	3144905	1088331	10
8	A002	R051	02-00-00	05-02-11	08:00:00	REGULAR	3144941	1088420	36
9	A002	R051	02-00-00	05-02-11	12:00:00	REGULAR	3145094	1088753	153
10	A002	R051	02-00-00	05-02-11	16:00:00	REGULAR	3145337	1088823	243
...									

#your code here

```
df['ENTRIES_hourly'] = df['ENTRIES'] - df['ENTRIES'].shift(periods=1)
df = df.fillna(1)
```

return df

Good job! Your code worked perfectly. Your output below:

Sn	Unnamed: 0	C/A	UNIT	SCP	DATE	TIME	DESC	ENTRIES	EXIT
0	ENTRIES_hourly								
51		0	A002	R051	02-00-00	05-01-11	00:00:00	REGULAR	3144312
		1							10881

9 - Get Hourly Exits:

import pandas

def get_hourly_exits(df):

'''

The data in the MTA Subway Turnstile data reports on the cumulative number of entries and exits per row. Assume that you have a dataframe called df that contains only the rows for a particular turnstile machine (i.e., unique SCP, C/A, and UNIT). This function should change these cumulative exit numbers to a count of exits since the last reading

(i.e., exits since the last row in the dataframe).

More specifically, you want to do two things:

- 1) Create a new column called EXITSn_hourly
- 2) Assign to the column the difference between EXITSn of the current row and the previous row. If there is any NaN, fill/replace it with 0.

You may find the pandas functions `shift()` and `fillna()` to be helpful in this exercise.

Example dataframe below:

```
Unnamed: 0  C/A  UNIT  SCP  DATEn  TIMEn  DESCn  ENTRIESn  EXITSn  ENTRIESn_hourly
EXITSn_hourly
0      0  A002  R051  02-00-00  05-01-11  00:00:00  REGULAR  3144312  1088151      0      0
1      1  A002  R051  02-00-00  05-01-11  04:00:00  REGULAR  3144335  1088159     23      8
2      2  A002  R051  02-00-00  05-01-11  08:00:00  REGULAR  3144353  1088177     18     18
3      3  A002  R051  02-00-00  05-01-11  12:00:00  REGULAR  3144424  1088231     71     54
4      4  A002  R051  02-00-00  05-01-11  16:00:00  REGULAR  3144594  1088275    170     44
5      5  A002  R051  02-00-00  05-01-11  20:00:00  REGULAR  3144808  1088317    214     42
6      6  A002  R051  02-00-00  05-02-11  00:00:00  REGULAR  3144895  1088328     87     11
7      7  A002  R051  02-00-00  05-02-11  04:00:00  REGULAR  3144905  1088331     10      3
8      8  A002  R051  02-00-00  05-02-11  08:00:00  REGULAR  3144941  1088420     36     89
9      9  A002  R051  02-00-00  05-02-11  12:00:00  REGULAR  3145094  1088753    153    333
'''
```

#your code here

```
df['EXITSn_hourly'] = df['EXITSn'] - df['EXITSn'].shift(periods=1)
df = df.fillna(0)
return df
```

Good job! Your code worked perfectly. Your output below:

```
Unnamed: 0  C/A  UNIT  SCP  DATEn  TIMEn  DESCn  ENTRIESn  EXIT
Sn  ENTRIESn_hourly  EXITSn_hourly
0      0  A002  R051  02-00-00  05-01-11  00:00:00  REGULAR  3144312  10881
51      0      0
```

```
import pandas
import time as ti
```

III

Write a function to extract the hour part from the input variable `time` and return it as an integer. For example:

- Please return hour as an integer.

III

```
hour = ti.strptime(time, "%H:%M:%S")[3]
return hour
```

	Unnamed: 0	UNIT	DATEn	TIMEn	DESCn	ENTRIESn_hourly	EXITSn_hourly	H
0	0	R022	05-01-11	00:00:00	REGULAR	0	0	
1	1	R022	05-01-11	04:00:00	REGULAR	562	173	

11 - Reformat Subway Dates:

```
import datetime
```

```
import time as timer
```

```
def reformat_subway_dates(date):
```

```
    """
```

The dates in our subway data are formatted in the format month-day-year.

The dates in our weather underground data are formatted year-month-day.

In order to join these two data sets together, we'll want the dates formatted the same way. Write a function that takes as its input a date in the MTA Subway data format, and returns a date in the weather underground format.

Hint:

There is a useful function in the datetime library called `strptime`.

More info can be seen here:

<http://docs.python.org/2/library/datetime.html#datetime.datetime.strptime>

```
    """
```

```
    # your code here
```

```
    temp = timer.strptime(date, "%m-%d-%y")
```

```
    t1 = temp[0]
```

```
    t2 = temp[1]
```

```
    t3 = temp[2]
```

```
    dt = datetime.datetime(temp[0], temp[1], temp[2])
```

```
    date_formatted = dt.strftime("%Y-%m-%d")
```

```
    return date_formatted
```

Good job! Your code worked perfectly. Your output below:

Hour	Unnamed: 0	UNIT	DATEn	TIMEn	DESCn	ENTRIESn_hourly	EXITSn_hourly
0	0	R022	2011-05-01	00:00:00	REGULAR	0	0
1	1	R022	2011-05-01	04:00:00	REGULAR	562	173
4							